



Hewlett Packard
Enterprise

Malicious website detection with large scale belief propagation on Spark

Mijung Kim, Jun Li, Manish Marwah, Alexander Ulanov,
Carlos Zubieta

Hewlett Packard Labs


Outline

- Introduction
- (Loopy) Belief Propagation
- Implementation on Spark
- Experiments
- Conclusions and Future Work

Introduction

Probabilistic Graphical Models

- Combine **graphs** and **probability theory**
- **Vertices**: random variables
- **Edges**: relationships between variables
- Model joint probability distribution
- Used for probabilistic reasoning




Graphical Model: An Example

Markov Random Field Model


- Model parameters
 - Priors / node factors
 - Edge factors

- Queries
 - $P(B)$
 - $P(B|A, C)$
 - $P(E|C)$
 -




Applications

- Computer vision
- Malware detection
- Fraud detection
- Bioinformatics
- Recommender system
-



Inference in Graphical Models

- Queries of the graphical model
 - $P(A \mid B, C) ?$
 - $P(D) ?$
 - *Inference*
- Methods
 - Exact
 - Approximate
 - Variational methods
 - **Loopy Belief Propagation**
 - Sampling based (e.g., Gibbs sampling)



Loopy Belief Propagation Algorithm

- Message passing iterative algorithm
- Exact on trees, approximate on graphs with loops


- Initialize messages
- At each node
 - **Read** messages from neighbors
 - **Update** own marginal probability (belief)

$$b(x_i) = \phi(x_i) \prod_{j \in \mathcal{N}(i)} m_{j \rightarrow i}$$

- **Send** updated messages to all neighbors

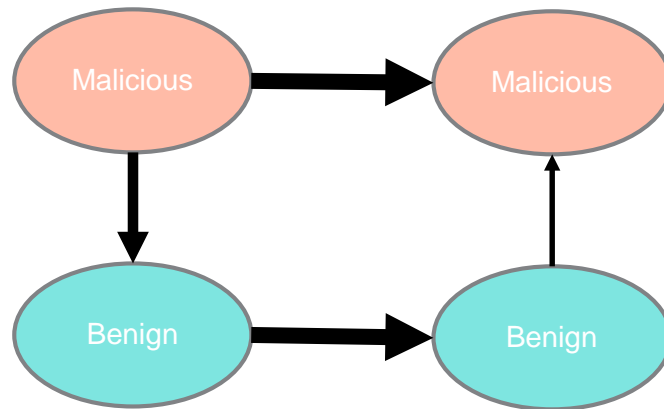
$$m_{i \rightarrow j}(x_j) = \sum_{x_i} \phi(x_i) \psi_{ij}(x_i, x_j) \prod_{k \in \mathcal{N}(i) \setminus j} m_{k \rightarrow i}$$

- Repeat until convergence




Application: Detection/Ranking of Malicious Websites

- Goal: To infer maliciousness score of a webhost/webpage
- Given web graph, exploit principle of “homophily”
- Malicious site → Malicious site
- Benign site → Benign site
-
- more likely than
- Benign site → Malicious site



Transform into a Graphical Model

- Model web graph as a Markov random field
 - Vertices: probability of a host being malicious
 - Edges: hyperlinks
- Exploit web link structure
- Use whitelist and blacklist to assign priors for some nodes
- Use uniform prior for rest
- Edge factors are assigned based on domain knowledge
- Perform inference (run BP) to estimate marginal probability
- Estimate score and rank



Number/type of nodes not representative of real graph

Challenges of running BP on Large-Scale Graphical Models

- Large Size/Random access
 - Use large memory machine to keep graph in memory
- Communication bound
 - Use shared memory
- Parallelization
 - Synchronous processing (BSP)
- Partitioning
 - Heuristic to minimize vertex replication

Implementation

Requirements

- Scalability
 - Apache Spark GraphX
- Generalized graph representation
 - Factor graph format
 - Handles factor of any order and variable domain of any size
- Numerical stability
 - Log domain calculations to delay underflow (overflow)
 - Factor math

Improvements

- Communication; memory



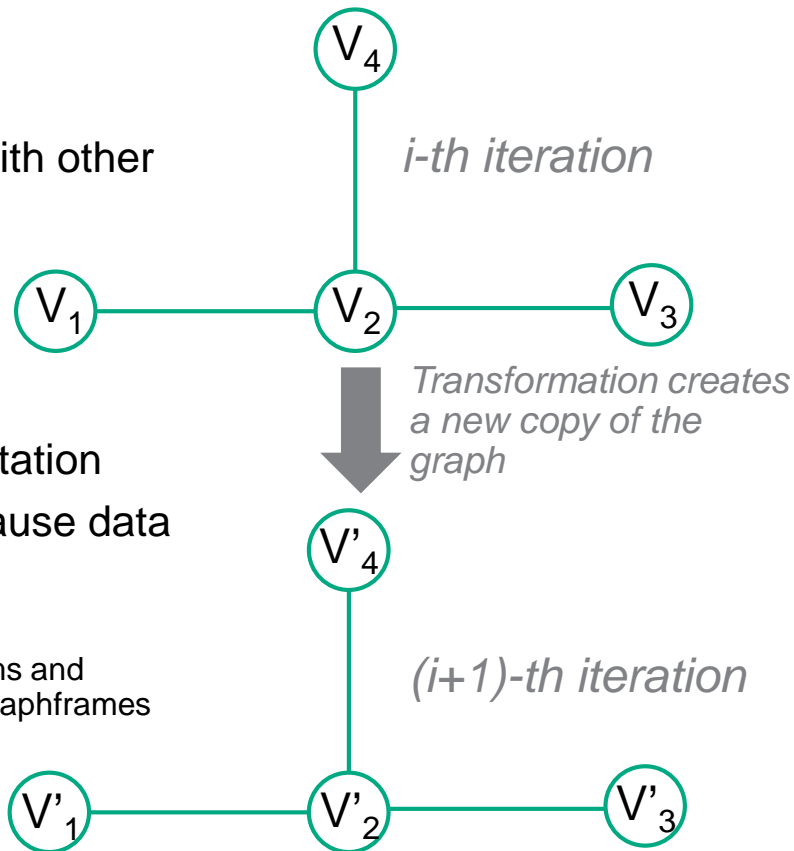
Scalability with Apache Spark GraphX

–Pros

- Provided by Apache Spark => easy to combine with other big data workloads
- Convenient message scatter/gather API

–Cons

- Is not actively developed
- Large memory overhead due to internal representation
- Requires 2X memory for iterative algorithms because data is immutable in Spark
 - Graph on previous iteration and graph on this iteration
 - Graphframes library does not properly support iterative algorithms and outsources them to GraphX. Will move our implementation to Graphframes once this issue is addressed



Graph representation in GraphX

- Graph in GraphX: VertexRDD, EdgeRDD, routing table
 - Vertex cut - replicated vertices
 - Routing table - data structure for updating replicated vertices
- Spark's computation model (BSP):


$$t = t_{cp} + t_{cm}$$

- Estimations for Belief Propagation

$$t_{cp} = \max_{i \in [1, n]} (E_i) / (F \cdot n) \cdot (S + 2 \cdot (S + S^2)) \quad t_{cm} = 32 / B \cdot r \cdot V \cdot S$$

– E - edges, V - vertices, S - number of states, F - FLOPS, B - bandwidth, r - replication factor


- Insight: communication is proportional to the number of vertices, computation is relatively inexpensive
- More details about the model: “Modeling Scalability of Distributed Machine Learning” *ICDE 2017*



Graph Representation for BP

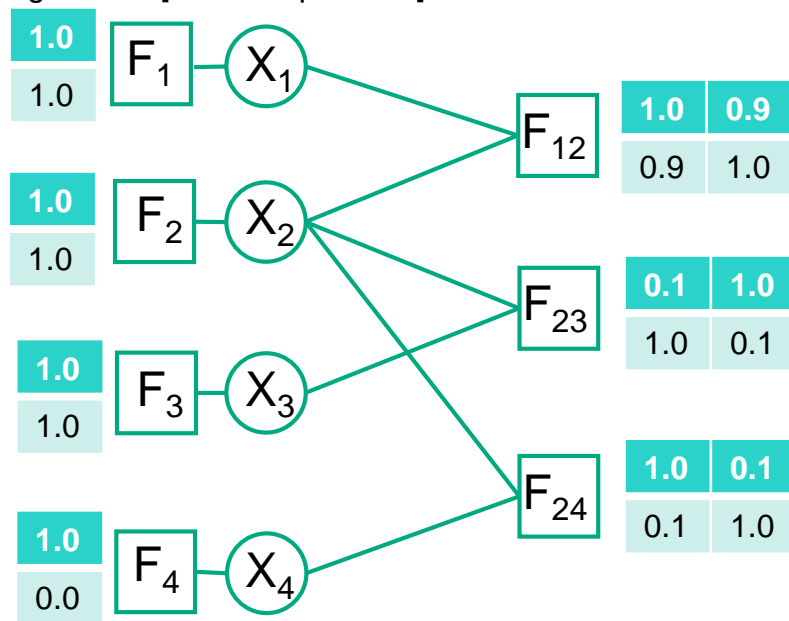
– Pairwise

- Variables are represented as vertices
- Factors are stored on edges
- Number of vertices == number of variables
- Number of edges == number of factors




– Factor graph (convenient in many domains)

- Can represent higher-order factors (can be converted to pairwise)
- Variables and factors are represented as vertices
- Priors are factors, but can be merged to variable vertex
- Number of vertices == number of variables + number of factors
- Number of edges == 2 [if factors pairwise] * the number of factors



Experiments: malicious web sites detection

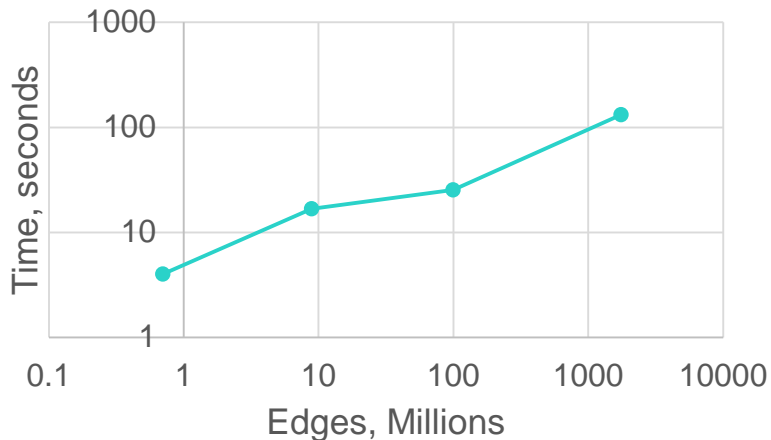
- Dataset
 - Host graph from commoncrawl/web-data-commons
 - 101.7M hosts, 1.75B edges
- Conversion for use-case
 - Host has two states: malicious and normal
 - Host priors from white and black lists
 - Whitelist Alexa500 (17.8M hosts in intersection)
 - Blacklist urlblacklist.com (0.93M hosts in intersection)
 - Factor represents link direction
- Goal: estimate probability of a host to be malicious (normal)
- Converged in 20 iterations, 250 seconds
- Result validation
 - Qualitatively on a subset
 - Malicious sites change quickly



- Hardware: SuperDome X
 - 16x Xeon E7-2890 v2 @ 2.80GHz
 - 11TB shared RAM
- BP algorithm on Spark 1.6.1
 - Shuffle dir in tempfs
 - 24 to 45 workers
 - 32 to 164GB RAM per worker
 - 1 core per worker

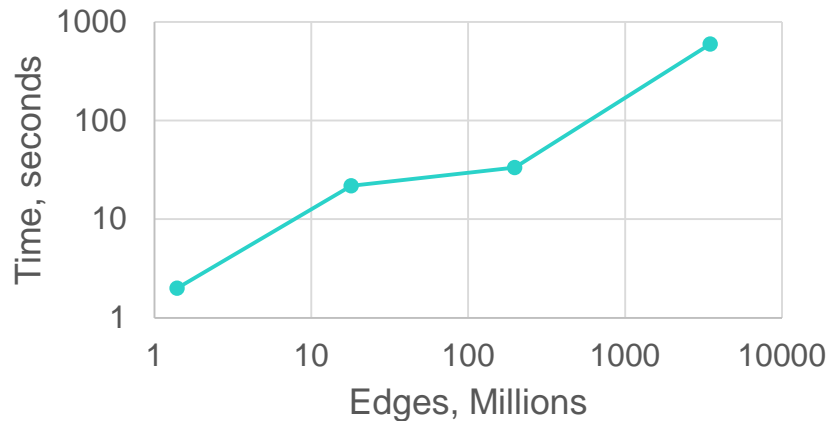
Experiments: various graphs

Pairwise iteration time



Vertices	Edges	Size on disk	Size in Spark	Iteration time, s.	Iterations
0.2M	0.7M	0.03GB	0.4GB	4	8
1.6M	8.9M	0.3GB	5.2GB	16.8	8
16.2M	99.3M	5.4GB	59.7GB	25.3	8
101.7M	1.75B	64GB	840.5GB	132	20

Factor Graph iteration time



Vertices	Edges	Size on disk	Size in Spark	Iteration time, s	Iterations
0.9M	1.4M	0.05GB	1.2GB	1.3	14
10.7M	18M	0.6GB	13GB	21.8	15
115.5M	198.5M	6.9GB	154.2GB	33.5	16
1.9B	3.5B	118G	2485G	600	40

Improvements

- Hewlett Packard Labs project: Spark for large memory machines (Spark4TM)
- Spark communication layer rewritten
 - Shared memory shuffle engine
- Off-heap memory store introduced
 - Mutable data structure
 - Breaks fault-tolerance model, however
 - The computation takes 10s of minutes at maximum; We can do checkpointing
- Clever graph partitioning
 - Based on heuristic to minimize vertex replication factor
- Spark4TM (custom shuffling + off heap + partitioning) provides 9 s. per iteration on 101M graph vs 132 s. in vanilla Spark
 - Code: <https://github.com/HewlettPackard/sparkle>

Project “Sandpiper”

– Open source implementation

– <https://github.com/HewlettPackard/sandpiper>

– <https://spark-packages.org/package/HewlettPackard/sandpiper>

File format: variables (id prior)

```
1 1.0 1.0
2 1.0 0.0
...
```

– Example of use

– BP for pairwise factors

– Two separate for files variables and factors

– Each record on a separate line

– Loading is automatically parallelized

File format: factors (id1 id2 factor table in column major format)

```
1 2 1.0 0.9 0.9 1.0
2 3 0.1 1.0 1.0 0.1
...
```

Scala

```
import sparkle.graph._
val graph = PairwiseBP.loadPairwiseGraph(sc, variableFile, factorFile)
val beliefs = PairwiseBP(graph, maxIterations = 50, epsilon = 1e-3)
```

Example of use: Factor Graph BP

File format

- BP for factor graph
 - libDAI file format with explicit factor ID
 - Factor ID enables splitting the file into parts for loading parallelism

```
# number of factors in the file
7

# factor id preceded with 3 hashes (unique, must not
intersect with variable name/id)
### 5
# number of vars
1
# name of vars
1
# number of values of vars
2
# number of non-zero entries in factor table
2
# non-zero factor table entries
0 1
1 1
```

Scala

```
import sparkle.graph._
val graph = Utils.loadLibDAIToFactorGraph(sc, inputPath)
val beliefs = BP(graph, maxIterations = 50, epsilon = 1e-3)
```

Conclusions and Future Work

- Belief propagation algorithm for large graphs
- Application of malicious site detection using BP
- Open source: project “sandpiper”
- Demo at HP Discover
- Future work
 - Re-compute only non-converged nodes
 - Incremental computations
 - Incremental graph construction



TOP 200 SITES

RANK	SCORE	LOCATION
1	1	Atlanta, United States
2	1	Newark, United States
3	1	Gostar, Iran
4	1	Boardman, United States
5	1	United States
6	1	Tempe, United States
7	1	San Antonio, United States
8	1	United Kingdom
9	1	United Kingdom
10	0.99997	St Louis, United States
11	0.99997	Atlanta, United States
12	0.99924	Netherlands
13	0.99912	United Kingdom
14	0.99427	St Louis, United States
15	0.99211	Netherlands
16	0.98736	United Kingdom
17	0.98439	Germany
18	0.98059	Germany
19	0.93884	St Louis, United States
20	0.91777	United Kingdom
21	0.91374	United Kingdom
22	0.90919	Turkey

#2

NN12.118 ***
Newark, United States
Score: 1

Spark for
The Machine

Hewlett Packard
Enterprise

Iteration 10:

Spark for
The Machine 0:13
Spark 2:46

Speedup 12.8x



14.8
Seconds

Data Filtering 18
Inference Processing 13



100 %
Convergence
of 101,717,775 sites

Start Over

View Summary

Created by
Hewlett Packard
Labs

Thank you



Spark for large pools of memory

- Make HPE large-memory servers accessible to customers and developers
- Can in-memory analytics perform better with big shared memory?
- Apache Spark as a platform



Superdome X
8 blades x3TB